

# **Software Engineering for Data Scientists**

## ***Manipulating Data with Python***

### **CSE 599 B1**

#### **Today's Objectives**

1. Opening & Navigating the IPython Notebook
2. Simple Math in the IPython Notebook
3. Loading data with pandas
4. Cleaning and Manipulating data with pandas
5. Visualizing data with pandas

## 1. Opening and Navigating the IPython Notebook

We will start today with the interactive environment that we will be using often through the course: the [IPython/Jupyter Notebook](http://ipython.org) (<http://ipython.org>).

We will walk through the following steps together:

1. Download [miniconda\(\)](#) (be sure to get Version 3.5) and install it on your system (hopefully you have done this before coming to class)
2. Use the conda command-line tool to update your package listing and install the IPython notebook:

Update conda's listing of packages for your system:

```
$ conda update conda
```

Install IPython notebook and all its requirements

```
$ conda install ipython-notebook
```

3. Navigate to the directory containing the course material. For example:

```
$ cd ~/courses/CSE599/
```

You should see a number of files in the directory, including these:

```
$ ls
...
Breakout-Simple-Math.ipynb
CSE599_Lecture_2.ipynb
...
```

4. Type `ipython notebook` in the terminal to start the notebook

```
$ ipython notebook
```

If everything has worked correctly, it should automatically launch your default browser

5. Click on `CSE599_Lecture_2.ipynb` to open the notebook containing the content for this lecture.

With that, you're set up to use the IPython notebook!

## 2. Simple Math in the IPython Notebook

Now that we have the IPython notebook up and running, we're going to do a short breakout exploring some of the mathematical functionality that Python offers.

Please open [Breakout-Simple-Math.ipynb](#) ([Breakout-Simple-Math.ipynb](#)), find a partner, and make your way through that notebook, typing and executing code along the way.

## 3. Loading data with pandas

With this simple Python computation experience under our belt, we can now move to doing some more interesting analysis.

## Python's Data Science Ecosystem

In addition to Python's built-in modules like the `math` module we explored above, there are also many often-used third-party modules that are core tools for doing data science with Python. Some of the most important ones are:

### **numpy** (<http://numpy.org/>): Numerical Python

Numpy is short for "Numerical Python", and contains tools for efficient manipulation of arrays of data. If you have used other computational tools like IDL or MatLab, Numpy should feel very familiar.

### **scipy** (<http://scipy.org/>): Scientific Python

Scipy is short for "Scientific Python", and contains a wide range of functionality for accomplishing common scientific tasks, such as optimization/minimization, numerical integration, interpolation, and much more. We will not look closely at Scipy today, but we will use its functionality later in the course.

### **pandas** (<http://pandas.pydata.org/>): Labeled Data Manipulation in Python

Pandas is short for "Panel Data", and contains tools for doing more advanced manipulation of labeled data in Python, in particular with a columnar data structure called a *Data Frame*. If you've used the [R](http://rstats.org) (<http://rstats.org>) statistical language (and in particular the so-called "Hadley Stack"), much of the functionality in Pandas should feel very familiar.

### **matplotlib** (<http://matplotlib.org/>): Visualization in Python

Matplotlib started out as a Matlab plotting clone in Python, and has grown from there in the 15 years since its creation. It is the most popular data visualization tool currently in the Python data world (though other recent packages are starting to encroach on its monopoly).

## Installing Pandas & friends

Because the above packages are not included in Python itself, you need to install them separately. While it is possible to install these from source (compiling the C and/or Fortran code that does the heavy lifting under the hood) it is much easier to use a package manager like `conda`. All it takes is to run

```
$ conda install numpy scipy pandas matplotlib
```

and (so long as your `conda` setup is working) the packages will be downloaded and installed on your system.

## Loading Data with Pandas

```
In [1]: import numpy
        numpy.__path__
```

```
Out[1]: ['/home/ubuntu/miniconda2/envs/python3/lib/python2.7/site-packages/numpy']
```

```
In [2]: import pandas
```

Because we'll use it so much, we often import under a shortened name using the `import ... as ...` pattern:

```
In [3]: import pandas as pd
```

Now we can use the `read_csv` command to read the comma-separated-value data:

```
In [4]: data = pd.read_csv('2015_trip_data.csv')
```

*Note: strings in Python can be defined either with double quotes or single quotes*

## Viewing Pandas Dataframes

The `head()` and `tail()` methods show us the first and last rows of the data

```
In [5]: data.head()
```

```
Out[5]:
```

	trip_id	starttime	stoptime	bikeid	tripduration	from_station_name	to_station_name	from
0	431	10/13/2014 10:31	10/13/2014 10:48	SEA00298	985.935	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
1	432	10/13/2014 10:32	10/13/2014 10:48	SEA00195	926.375	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
2	433	10/13/2014 10:33	10/13/2014 10:48	SEA00486	883.831	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
3	434	10/13/2014 10:34	10/13/2014 10:48	SEA00333	865.937	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
4	435	10/13/2014 10:34	10/13/2014 10:49	SEA00202	923.923	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI

In [6]: `data.tail()`

Out[6]:

	trip_id	starttime	stoptime	bikeid	tripduration	from_station_name	to_station_name
<b>142841</b>	156796	10/12/2015 20:41	10/12/2015 20:47	SEA00358	377.183	E Pine St & 16th Ave	Summit Ave & E Denny Way
<b>142842</b>	156797	10/12/2015 20:43	10/12/2015 20:48	SEA00399	303.330	Bellevue Ave & E Pine St	Summit Ave E & E Republican St
<b>142843</b>	156798	10/12/2015 21:03	10/12/2015 21:06	SEA00204	165.597	Harvard Ave & E Pine St	E Harrison St & Broadway Ave E
<b>142844</b>	156799	10/12/2015 21:35	10/12/2015 21:41	SEA00073	388.576	Pine St & 9th Ave	3rd Ave & Broad St
<b>142845</b>	156800	10/12/2015 22:45	10/12/2015 22:51	SEA00033	391.885	NE 42nd St & University Way NE	Eastlake Ave E & E Allison St

The shape attribute shows us the number of elements:

In [7]: `data.shape`

Out[7]: (142846, 12)

The columns attribute gives us the column names

In [8]: `data.columns`

Out[8]: Index([u'trip\_id', u'starttime', u'stoptime', u'bikeid', u'tripduration',  
u'from\_station\_name', u'to\_station\_name', u'from\_station\_id',  
u'to\_station\_id', u'usertype', u'gender', u'birthyear'],  
dtype='object')

The index attribute gives us the index names

In [9]: `data.index`

Out[9]: RangeIndex(start=0, stop=142846, step=1)

The dtypes attribute gives the data types of each column:

```
In [10]: data.dtypes
```

```
Out[10]: trip_id          int64  
starttime          object  
stoptime           object  
bikeid            object  
tripduration      float64  
from_station_name  object  
to_station_name    object  
from_station_id    object  
to_station_id      object  
usertype           object  
gender             object  
birthyear         float64  
dtype: object
```

## 4. Manipulating data with pandas

Here we'll cover some key features of manipulating data with pandas

Access columns by name using square-bracket indexing:

```
In [11]: data["usertype"]
```

```

Out[11]: 0          Annual Member
         1          Annual Member
         2          Annual Member
         3          Annual Member
         4          Annual Member
         5          Annual Member
         6          Annual Member
         7          Annual Member
         8          Annual Member
         9          Annual Member
        10          Annual Member
        11          Annual Member
        12          Annual Member
        13          Annual Member
        14          Annual Member
        15          Annual Member
        16          Annual Member
        17          Annual Member
        18          Annual Member
        19          Annual Member
        20          Annual Member
        21          Annual Member
        22          Annual Member
        23          Annual Member
        24          Annual Member
        25          Annual Member
        26          Annual Member
        27          Annual Member
        28          Annual Member
        29          Annual Member
        ...
    142816          Annual Member
    142817          Annual Member
    142818          Annual Member
    142819          Annual Member
    142820      Short-Term Pass Holder
    142821      Short-Term Pass Holder
    142822      Short-Term Pass Holder
    142823      Short-Term Pass Holder
    142824          Annual Member
    142825          Annual Member
    142826          Annual Member
    142827          Annual Member
    142828          Annual Member
    142829          Annual Member
    142830          Annual Member
    142831          Annual Member
    142832          Annual Member
    142833      Short-Term Pass Holder
    142834          Annual Member
    142835          Annual Member
    142836          Annual Member
    142837          Annual Member
    142838          Annual Member
    142839          Annual Member
    142840          Annual Member
    142841          Annual Member
    142842          Annual Member
    142843          Annual Member
    142844      Short-Term Pass Holder
    142845          Annual Member
Name: usertype, dtype: object

```



Mathematical operations on columns happen *element-wise*:

```
In [12]: data['tripduration'] / 60
```

```
Out[12]: 0      16.432250
          1      15.439583
          2      14.730517
          3      14.432283
          4      15.398717
          5      13.480083
          6       9.945250
          7       9.868850
          8       9.772450
          9       9.793900
         10       9.414983
         11      10.335683
         12      10.568117
         13      10.238933
         14      10.024383
         15      10.313017
         16      10.284750
         17      10.000833
         18       8.328900
         19       9.588450
         20       9.530117
         21       9.396050
         22       7.294817
         23       8.052567
         24       7.989700
         25       8.892200
         26       8.027150
         27       7.679533
         28       7.652750
         29      11.340950
          ...
142816      8.671450
142817     14.607983
142818      4.938550
142819     11.882017
142820     21.947550
142821     21.883067
142822     21.240667
142823     18.747133
142824     22.390117
142825     12.643100
142826      4.613267
142827      6.036600
142828      4.168183
142829      4.278083
142830      6.128367
142831      7.455983
142832      2.109933
142833      9.045133
142834      4.308700
142835     15.671783
142836     10.276817
142837     10.726967
142838      9.055200
142839      3.375017
142840      8.962267
142841      6.286383
142842      5.055500
142843      2.759950
142844      6.476267
142845      6.531417
Name: tripduration, dtype: float64
```

Columns can be created (or overwritten) with the assignment operator. Let's create a *tripminutes* column with the number of minutes for each trip

```
In [13]: data['tripminutes'] = data['tripduration'] / 60
```

```
In [14]: data.head()
```

```
Out[14]:
```

	trip_id	starttime	stoptime	bikeid	tripduration	from_station_name	to_station_name	froi
0	431	10/13/2014 10:31	10/13/2014 10:48	SEA00298	985.935	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
1	432	10/13/2014 10:32	10/13/2014 10:48	SEA00195	926.375	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
2	433	10/13/2014 10:33	10/13/2014 10:48	SEA00486	883.831	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
3	434	10/13/2014 10:34	10/13/2014 10:48	SEA00333	865.937	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI
4	435	10/13/2014 10:34	10/13/2014 10:49	SEA00202	923.923	2nd Ave & Spring St	Occidental Park / Occidental Ave S & S Washing...	CBI

## Working with Times

One trick to know when working with columns of times is that Pandas `DatetimeIndex` provides a nice interface for working with columns of times:

```
In [15]: times = pd.DatetimeIndex(data['starttime'])
```

With it, we can extract, the hour of the day, the day of the week, the month, and a wide range of other views of the time:

```
In [16]: times
```

```
Out[16]: DatetimeIndex(['2014-10-13 10:31:00', '2014-10-13 10:32:00',
                        '2014-10-13 10:33:00', '2014-10-13 10:34:00',
                        '2014-10-13 10:34:00', '2014-10-13 10:34:00',
                        '2014-10-13 11:35:00', '2014-10-13 11:35:00',
                        '2014-10-13 11:35:00', '2014-10-13 11:35:00',
                        ...
                        '2015-10-12 20:09:00', '2015-10-12 20:11:00',
                        '2015-10-12 20:18:00', '2015-10-12 20:39:00',
                        '2015-10-12 20:41:00', '2015-10-12 20:41:00',
                        '2015-10-12 20:43:00', '2015-10-12 21:03:00',
                        '2015-10-12 21:35:00', '2015-10-12 22:45:00'],
                        dtype='datetime64[ns]', length=142846, freq=None)
```

```
In [17]: times.dayofweek
```

```
Out[17]: array([0, 0, 0, ..., 0, 0, 0], dtype=int32)
```

```
In [18]: times.month
```

```
Out[18]: array([10, 10, 10, ..., 10, 10, 10], dtype=int32)
```

*Note: math functionality can be applied to columns using the NumPy package: for example:*

```
In [19]: import numpy as np  
         np.exp(data['tripminutes'])
```

```
Out[19]: 0      1.369101e+07
          1      5.073712e+06
          2      2.496791e+06
          3      1.852939e+06
          4      4.870546e+06
          5      7.150325e+05
          6      2.085294e+04
          7      1.931911e+04
          8      1.754370e+04
          9      1.792407e+04
         10      1.227087e+04
         11      3.081273e+04
         12      3.887539e+04
         13      2.797127e+04
         14      2.257015e+04
         15      3.012217e+04
         16      2.928264e+04
         17      2.204483e+04
         18      4.141859e+03
         19      1.459523e+04
         20      1.376820e+04
         21      1.204073e+04
         22      1.472647e+03
         23      3.141849e+03
         24      2.950412e+03
         25      7.275007e+03
         26      3.063000e+03
         27      2.163610e+03
         28      2.106430e+03
         29      8.419998e+04
          ...
        142816    5.833952e+03
        142817    2.208852e+06
        142818    1.395677e+02
        142819    1.446420e+05
        142820    3.401730e+09
        142821    3.189298e+09
        142822    1.677661e+09
        142823    1.386043e+08
        142824    5.295465e+09
        142825    3.096197e+05
        142826    1.008129e+02
        142827    4.184678e+02
        142828    6.459799e+01
        142829    7.210211e+01
        142830    4.586864e+02
        142831    1.730185e+03
        142832    8.247691e+00
        142833    8.477182e+03
        142834    7.434378e+01
        142835    6.399839e+06
        142836    2.905125e+04
        142837    4.556826e+04
        142838    8.562950e+03
        142839    2.922477e+01
        142840    7.803024e+03
        142841    5.372069e+02
        142842    1.568830e+02
        142843    1.579905e+01
        142844    6.495415e+02
        142845    6.863699e+02
        Name: tripminutes, dtype: float64
```

## Simple Grouping of Data

The real power of Pandas comes in its tools for grouping and aggregating data. Here we'll look at *value counts* and the basics of *group-by* operations.

### Value Counts

Pandas includes an array of useful functionality for manipulating and analyzing tabular data. We'll take a look at two of these here.

The `pandas.value_counts` returns statistics on the unique values within each column.

We can use it, for example, to break down rides by gender:

```
In [20]: pd.value_counts(data['gender'])
```

```
Out[20]: Male      67608  
         Female    18245  
         Other      1507  
         Name: gender, dtype: int64
```

Or to break down rides by age:



```
In [21]: pd.value_counts(data['birthyear']).sort_index()
```

```
Out[21]: 1936.0      6
          1939.0     23
          1942.0      2
          1943.0     11
          1944.0      1
          1945.0    115
          1946.0     39
          1947.0    244
          1948.0     34
          1949.0    154
          1950.0    657
          1951.0    251
          1952.0    204
          1953.0    337
          1954.0    152
          1955.0    397
          1956.0    481
          1957.0    224
          1958.0    160
          1959.0    696
          1960.0    429
          1961.0    875
          1962.0   1769
          1963.0    828
          1964.0   1374
          1965.0   1510
          1966.0    761
          1967.0   1354
          1968.0   1085
          1969.0   1185
          1970.0   1056
          1971.0   1279
          1972.0   1921
          1973.0   1076
          1974.0   1497
          1975.0   1969
          1976.0   1577
          1977.0   2465
          1978.0   2063
          1979.0   1976
          1980.0   2236
          1981.0   4779
          1982.0   4629
          1983.0   3965
          1984.0   3815
          1985.0   5370
          1986.0   3492
          1987.0   9320
          1988.0   4188
          1989.0   2755
          1990.0   3605
          1991.0   2912
          1992.0   1798
          1993.0   1126
          1994.0    549
          1995.0    412
          1996.0    121
          1997.0     21
          1998.0     25
          1999.0      5
          Name: birthyear, dtype: int64
```

What else might we break down rides by?

```
In [22]: pd.value_counts(times.dayofweek)
```

```
Out[22]: 3    21505
         0    21266
         4    21097
         2    20748
         1    20465
         5    20358
         6    17407
         dtype: int64
```

*We can sort by the index rather than the counts if we wish:*

```
In [23]: pd.value_counts(times.dayofweek, sort=False)
```

```
Out[23]: 0    21266
         1    20465
         2    20748
         3    21505
         4    21097
         5    20358
         6    17407
         dtype: int64
```

```
In [24]: pd.value_counts(times.month)
```

```
Out[24]: 7    18808
         8    17046
         6    15999
         5    15548
         9    13134
         4    12898
        10    11081
         3     9980
        11     7823
         1     7368
         2     7330
        12     5831
         dtype: int64
```

```
In [25]: pd.value_counts(times.month, sort=False)
```

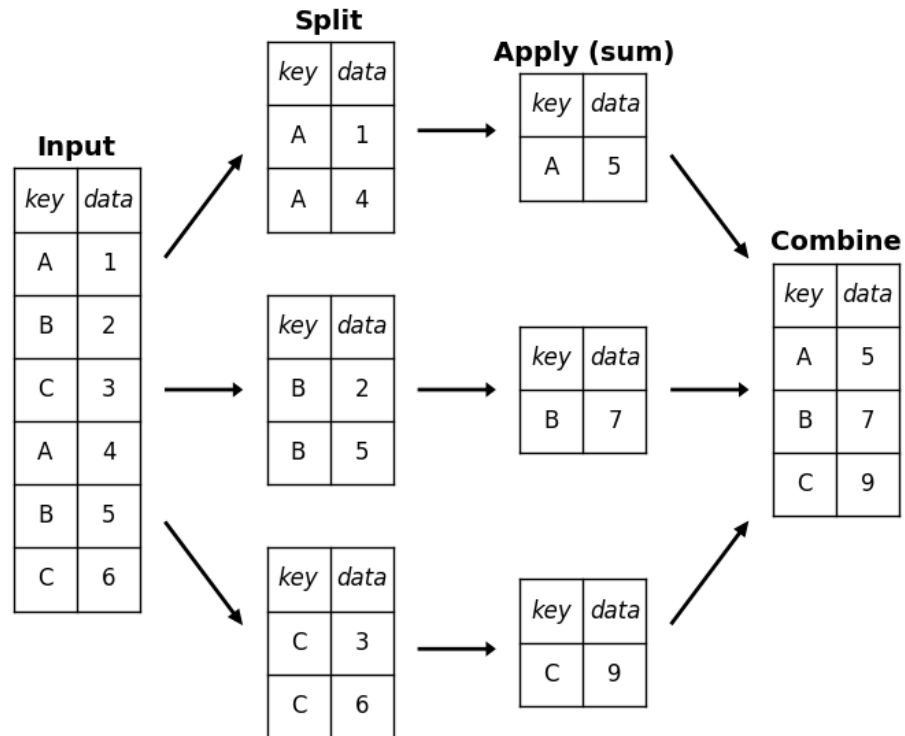
```
Out[25]: 1     7368
         2     7330
         3     9980
         4    12898
         5    15548
         6    15999
         7    18808
         8    17046
         9    13134
        10    11081
        11     7823
        12     5831
         dtype: int64
```

## Group-by Operation

One of the killer features of the Pandas dataframe is the ability to do group-by operations. You can visualize the group-by like this (image borrowed from the [Python Data Science Handbook](http://shop.oreilly.com/product/0636920034919.do) (<http://shop.oreilly.com/product/0636920034919.do>))

```
In [26]: from IPython.display import Image  
         Image('split_apply_combine.png')
```

Out[26]:



Let's break take this in smaller steps. First, let's look at the data by hour across all days in the year.

```
In [27]: pd.value_counts(times.hour)
```

```
Out[27]: 17    14163
         16    11629
         8     10967
         18    10382
         15     9850
         9     9751
         13     9575
         12     9571
         14     9096
         11     8864
         10     7761
         19     6939
         7     6093
         20     4792
         21     3730
         22     2484
         6     1855
         23     1749
         0     1022
         5       905
         1       682
         2       478
         4       316
         3       192
         dtype: int64
```

groupby allows us to look at the number of values for each column and each value.

In [28]: `data.groupby(times.hour).count()`

Out[28]:

	trip_id	starttime	stoptime	bikeid	tripduration	from_station_name	to_station_name	from_stat
0	1022	1022	1022	1022	1022	1022	1022	1022
1	682	682	682	682	682	682	682	682
2	478	478	478	478	478	478	478	478
3	192	192	192	192	192	192	192	192
4	316	316	316	316	316	316	316	316
5	905	905	905	905	905	905	905	905
6	1855	1855	1855	1855	1855	1855	1855	1855
7	6093	6093	6093	6093	6093	6093	6093	6093
8	10967	10967	10967	10967	10967	10967	10967	10967
9	9751	9751	9751	9751	9751	9751	9751	9751
10	7761	7761	7761	7761	7761	7761	7761	7761
11	8864	8864	8864	8864	8864	8864	8864	8864
12	9571	9571	9571	9571	9571	9571	9571	9571
13	9575	9575	9575	9575	9575	9575	9575	9575
14	9096	9096	9096	9096	9096	9096	9096	9096
15	9850	9850	9850	9850	9850	9850	9850	9850
16	11629	11629	11629	11629	11629	11629	11629	11629
17	14163	14163	14163	14163	14163	14163	14163	14163
18	10382	10382	10382	10382	10382	10382	10382	10382
19	6939	6939	6939	6939	6939	6939	6939	6939
20	4792	4792	4792	4792	4792	4792	4792	4792
21	3730	3730	3730	3730	3730	3730	3730	3730
22	2484	2484	2484	2484	2484	2484	2484	2484
23	1749	1749	1749	1749	1749	1749	1749	1749

Now, let's find the average length of a ride as a function of time of day:

```
In [29]: data.groupby(times.hour)['tripminutes'].mean()
```

```
Out[29]: 0      18.293162
         1      16.812000
         2      26.467510
         3      22.643443
         4      18.595762
         5      13.565035
         6      12.091993
         7      12.378344
         8      12.544350
         9      15.175861
        10      23.558911
        11      25.645489
        12      26.052903
        13      27.878785
        14      28.354453
        15      26.164124
        16      21.257375
        17      17.388788
        18      16.706635
        19      16.886609
        20      17.463562
        21      15.227905
        22      15.931296
        23      16.006255
        Name: tripminutes, dtype: float64
```

You can specify a groupby using the names of table columns and compute other functions, such as the mean.

```
In [30]: data.groupby(['gender'])['tripminutes'].mean()
```

```
Out[30]: gender
         Female      12.137525
         Male       9.547313
         Other      10.898911
        Name: tripminutes, dtype: float64
```

The simplest version of a groupby looks like this, and you can use almost any aggregation function you wish (mean, median, sum, minimum, maximum, standard deviation, count, etc.)

```
<data object>.groupby(<grouping values>).<aggregate>()
```

You can even group by multiple values: for example we can look at the trip duration by time of day and by gender:

```
In [31]: grouped = data.groupby([times.hour, 'gender'])['tripminutes'].mean()  
grouped
```



```
Out[31]: gender
0 Female 9.016608
  Male 8.750898
  Other 8.123788
1 Female 8.721844
  Male 7.865786
  Other 13.765683
2 Female 8.371583
  Male 8.880909
  Other 7.344022
3 Female 16.796675
  Male 8.087312
  Other 5.643350
4 Female 9.123506
  Male 12.134917
  Other 2.810767
5 Female 12.970906
  Male 8.261415
6 Female 10.954091
  Male 8.038184
  Other 9.752973
7 Female 12.392548
  Male 9.461130
  Other 10.525419
8 Female 11.309898
  Male 9.219013
  Other 10.129815
9 Female 10.881655
  Male 9.019535
  Other 11.167447
10 Female 12.152443
    ...
14 Female 13.085510
  Male 10.097013
  Other 11.780128
15 Female 12.521131
  Male 10.586964
  Other 10.580798
16 Female 12.614844
  Male 10.534829
  Other 13.845402
17 Female 12.927452
  Male 10.507247
  Other 11.990486
18 Female 12.414073
  Male 9.432061
  Other 9.477819
19 Female 12.008509
  Male 9.068067
  Other 9.579921
20 Female 11.716320
  Male 9.253255
  Other 8.051081
21 Female 11.241531
  Male 8.859545
  Other 7.606792
22 Female 11.088358
  Male 8.733132
  Other 6.618971
23 Female 11.032336
  Male 8.085837
  Other 5.183287
Name: tripminutes, dtype: float64
```

The `unstack()` operation can help make sense of this type of multiply-grouped data. What this technically does is split a multiple-valued index into an index plus columns:

In [32]: `grouped.unstack()`

Out[32]:

gender	Female	Male	Other
0	9.016608	8.750898	8.123788
1	8.721844	7.865786	13.765683
2	8.371583	8.880909	7.344022
3	16.796675	8.087312	5.643350
4	9.123506	12.134917	2.810767
5	12.970906	8.261415	NaN
6	10.954091	8.038184	9.752973
7	12.392548	9.461130	10.525419
8	11.309898	9.219013	10.129815
9	10.881655	9.019535	11.167447
10	12.152443	9.275972	11.615609
11	11.763075	9.066147	13.589810
12	12.436949	9.859569	10.503408
13	13.031992	9.675181	10.693493
14	13.085510	10.097013	11.780128
15	12.521131	10.586964	10.580798
16	12.614844	10.534829	13.845402
17	12.927452	10.507247	11.990486
18	12.414073	9.432061	9.477819
19	12.008509	9.068067	9.579921
20	11.716320	9.253255	8.051081
21	11.241531	8.859545	7.606792
22	11.088358	8.733132	6.618971
23	11.032336	8.085837	5.183287

## 5. Visualizing data with pandas

Of course, looking at tables of data is not very intuitive. Fortunately Pandas has many useful plotting functions built-in, all of which make use of the `matplotlib` library to generate plots.


Whenever you do plotting in the IPython notebook, you will want to first run this *magic command* which configures the notebook to work well with plots:

```
In [33]: %matplotlib inline
```

Now we can simply call the `plot()` method of any series or dataframe to get a reasonable view of the data:

```
In [34]: data.groupby([times.hour, 'usertype'])['tripminutes'].mean().unstack().plot()
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x7f56dc760c50>
```



## Adjusting the Plot Style

The default formatting is not very nice; I often make use of the [Seaborn](http://stanford.edu/~mwaskom/software/seaborn/) (<http://stanford.edu/~mwaskom/software/seaborn/>) library for better plotting defaults.

You should do this in bash

```
$ conda install seaborn
```

Then this in python

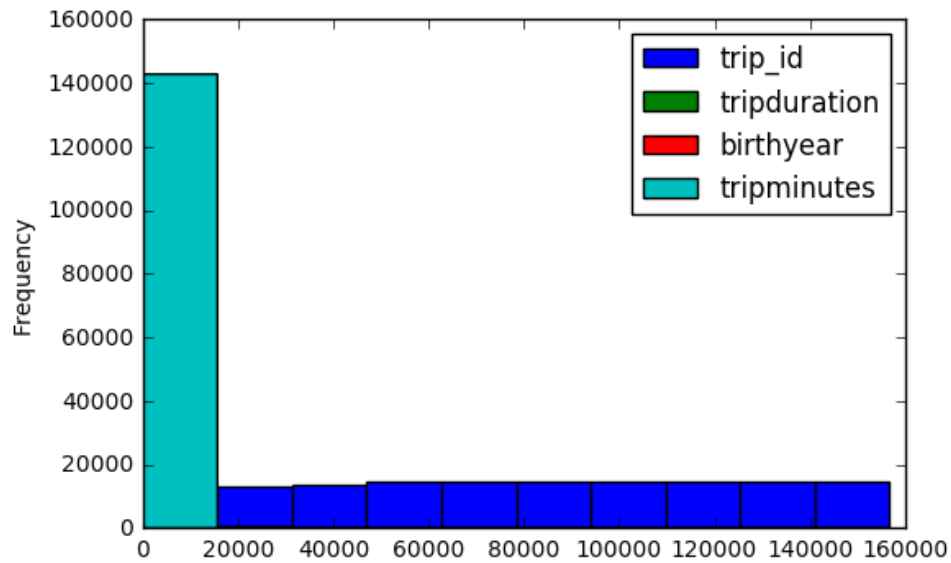
```
import seaborn
seaborn.set()
data.groupby([times.hour, 'usertype'])['tripminutes'].mean().unstack().plot()
```

## Other plot types

Pandas supports a range of other plotting types; you can find these by using the autocomplete on the `plot` method:

```
In [35]: data.plot.hist()
```


```
Out[35]: <matplotlib.axes._subplots.AxesSubplot at 0x7f56dc3f0290>
```



For example, we can create a histogram of trip durations:

```
In [36]: data['tripminutes'].plot.hist(bins=100)
```


```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7f56cd03b190>
```



If you'd like to adjust the x and y limits of the plot, you can use the `set_xlim()` and `set_ylim()` method of the resulting object:

```
In [37]: plot = data['tripminutes'].plot.hist(bins=500)
plot.set_xlim(0, 50)
```

Out[37]: (0, 50)



## Breakout: Exploring the Data

1. Make a plot of the total number of rides as a function of month of the year (You'll need to extract the month, use a groupby, and find the appropriate aggregation to count the number in each group).

In [ ]:

1. Split this plot by gender. Do you see any seasonal ridership patterns by gender?

In [ ]:

1. Split this plot by user type. Do you see any seasonal ridership patterns by usertype?

In [ ]:

1. Repeat the above three steps, counting the number of rides by time of day rather than by month.

In [ ]:

1. Are there any other interesting insights you can discover in the data using these tools?

In [ ]:

In [ ]:

## Looking Forward to Homework

In the homework this week, you will have a chance to apply some of these patterns to a brand new (but closely related) dataset.