

```
In [1]: %load_ext sql
        %sql postgresql://dssg_student:password@sed-sql.csysa4zsfb6y4.us-east-1.rds.amazonaws.com/dssg2016

Out[1]: 'Connected: dssg_student@dssg2016'
```

Introduction to Databases with SQL

Valentina Staneva

eScience Institute

vms16@uw.edu

11/30/2016

Plan for Today:

- quickly overview functionality of databases and SQL
- learn fundamental SQL commands by exploring a Seattle Crime Dataset

Data source: <http://www.seattle.gov/seattle-police-department/crime-data/spd-data-sets>
(<http://www.seattle.gov/seattle-police-department/crime-data/spd-data-sets>)

Lesson source: https://github.com/valentina-s/SQL_tutorial/ (https://github.com/valentina-s/SQL_tutorial/)


Why Databases???

- data you need is in a public database
- you want to make your data easily accessible
- you can manage the security permissions of the data
- computational load is on the server side
- you cannot load data into RAM
- certain operations are optimized in a database
- you can implement efficient scalable algorithms in a database

What is a Database?

"A database system is basically a computerized recordkeeping system -- that is, a system whose overall purpose is to maintain information and to make that information available on demand." (Date 1986)

- conventional databases use the *relational* data model



Relational Data Model:

- data are structured into row/column format

crimesID	Offense type	Offense code	Date	Location
1	tresspass	5700	2015-01-28 09:30:00	12XX Block of E Pike St
2	larceny-theft	2300	2015-02-21 08:24:21	15XX Block of Aurora St

- each record has a unique identifier (primary key)

What is SQL?

Structured Query Language

- Domain-Specific Programming Language
- not Turing Complete (cannot build a robot which can follow any instruction :())

But great for data manipulation!!!

And human-readable!!!

Creating a Table:

```
CREATE TABLE seattlecrimesincidents
("crimesID" int,
 "Offense type" character,
 "Offense code" int,
 "Date" timestamp,
 "Location" character);
```

crimesID	Offense type	Offense code	Date	Location

- populating the database records:

```
INSERT INTO seattlecrimeincidents VALUES
```

```
(1,'trespass', 5700,'2015-01-28 09:30:00','12XX Block of E Pike St'),
```

```
(2,'larceny-theft',2300, '2015-02-21 08:24:21','15XX Block of Aurora St');
```

crimesID	Offense type	Offense code	Date	Location
1	tresspass	5700	2015-01-28 09:30:00	12XX Block of E Pike St
2	larceny-theft	2300	2015-02-21 08:24:21	15XX Block of Aurora St

Data in each column must be of the same type

Some common data types (<https://www.postgresql.org/docs/9.4/static/datatype.html>):

Name	Aliases	Description
boolean	bool	logical Boolean (true/false)
character [(n)]	char [(n)]	fixed-length character string
date		calendar date (year, month, day)
double precision	float8	double precision floating-point number (8 bytes)
integer	int, int4	signed four-byte integer
json		JSON data
money		currency amount
timestamp [(p)] [without time zone]		date and time (no time zone)
xml		XML data

NULL values

- missing data are a common feature of many datasets
- here the code for "tresspass" is not known so the data entry is "X"

crimesID	Offense type	Offense code	Date	Location
1	tresspass	X	2015-01-28 09:30:00	12XX Block of E Pike St
2	burglary	5710	2015-01-28 09:30:00	12XX Block of E Pike St
3	larceny-theft	2300	2015-02-21 08:24:21	15XX Block of Aurora St

NULL values

- conventionally, some value is used to represent missing data (e.g. "X" or -9999)
- relational databases introduced NULL values:
 - NULL is a state representing a lack of a value
 - NULL is not the same as zero!
 - NULL values are ignored in SELECT statements

Normalization (https://en.wikipedia.org/wiki/Database_normalization)

- minimize redundancy

Example: multiple offenses at the same time

crimesID	Offense type	Offense code	Date	Location
1	tresspass and burglary	5700 and 5710	2015-01-28 09:30:00	12XX Block of E Pike St
2	larceny-theft	2300	2015-02-21 08:24:21	15XX Block of Aurora St

INCORRECT: database will have problems searching these columns

Solution: create another row

crimesID	Offense type	Offense code	Date	Location
1	tresspass	5700	2015-01-28 09:30:00	12XX Block of E Pike St
2	burglary	5710	2015-01-28 09:30:00	12XX Block of E Pike St
3	larceny-theft	2300	2015-02-21 08:24:21	15XX Block of Aurora St

Selecting Rows:

```
SELECT *  
  FROM seattlecrimeincidents  
 WHERE "Offense code" = 5700;
```

- use a "WHERE" clause to select specific rows

crimesID	Offense type	Offense code	Date	Location
1	tresspass	5700	2015-01-28 09:30:00	12XX Block of E Pike St

Selecting Columns:

```
SELECT "Offense type", "Date"
FROM seattlecrimeincidents;
```

- use a comma separated list to select specific columns

Offense type	Date
tresspass	2015-01-28 09:30:00
larceny-theft	2015-02-21 08:24:21

Elementwise Functions on Columns

Example:

- use a function to extract a subset of a date (e.g. year, hour) from a column with type = "timestamp"

```
In [4]: %%sql
SELECT "Date Reported", date_part('hour', "Date Reported")
FROM seattlecrimeincidents
LIMIT 5;
```

5 rows affected.

Out[4]:

Date Reported	date_part
2015-01-26 13:25:00	13.0
2015-01-29 14:32:00	14.0
2015-01-22 04:35:00	4.0
2015-01-17 01:21:00	1.0
2015-02-02 06:48:00	6.0

Aggregate Functions on Columns

- examples: SUM(), MAX(), MIN(), AVG(), COUNT(), STDDEV()

Data Analysis:

- databases have powerful methods for analyzing data
- one of the most common tasks: applying statistics across groups
- to accomplish this we need to learn
 - how to GROUP sets of data
 - how to apply statistical functions to those groups

crimesID	Offense code	Date	Location	Damage
1	5700	2015-01-28 09:30:00	12XX Block of E Pike St	\$1,220
1	5700	2015-02-12 03:25:00	1XX Block of Aloha St	\$11,420
2	5710	2015-01-28 09:30:00	12XX Block of E Pike St	\$5,389
2	5710	2015-1-02 12:31:20	12XX Block of E Pine St	\$15,231
3	2300	2015-02-21 08:24:21	15XX Block of Aurora St	\$2,405

Q: What is the total damage that occurred for each offense type?

- data grouped by "Offense code":

crimesID	Offense code	Date	Location	Damage
1	5700	2015-01-28 09:30:00	12XX Block of E Pike St	\$1,220
2	5700	2015-02-12 03:25:00	1XX Block of Aloha St	\$11,420
3	5710	2015-01-28 09:30:00	12XX Block of E Pike St	\$5,389
4	5710	2015-1-02 12:31:20	12XX Block of E Pine St	\$15,231
5	2300	2015-02-21 08:24:21	15XX Block of Aurora St	\$2,405

- data grouped by "Offense code":

crimesID	Offense code	Date	Location	Damage
1	5700	2015-01-28 09:30:00	12XX Block of E Pike St	\$1,220
2	5700	2015-02-12 03:25:00	1XX Block of Aloha St	\$11,420
3	5710	2015-01-28 09:30:00	12XX Block of E Pike St	\$5,389
4	5710	2015-1-02 12:31:20	12XX Block of E Pine St	\$15,231
5	2300	2015-02-21 08:24:21	15XX Block of Aurora St	\$2,405

```
SELECT SUM("Damage")
  FROM seattlecrimeincidents
 GROUP BY "Offense code";
```

Offense code	totalDamage
5700	\$12,640
5710	\$20,620
2300	\$2,405

Column Aliasing:

- often we want to rename newly generated columns:

In [5]: `%%sql`
`SELECT "Date Reported", date_part('hour', "Date Reported") AS "reported`
`hour"`
`FROM seattlecrimeincidents`
`LIMIT 5;`

5 rows affected.

Out[5]:

Date Reported	reported hour
2015-01-26 13:25:00	13.0
2015-01-29 14:32:00	14.0
2015-01-22 04:35:00	4.0
2015-01-17 01:21:00	1.0
2015-02-02 06:48:00	6.0

Joining Tables

- well designed databases distribute data across multiple tables, for efficiency
- then we can JOIN data between tables as needed

crimesID	Offense code	Date	Location
1	5700	2015-01-28 09:30:00	12XX Block of E Pike St
1	5700	2015-02-12 03:25:00	1XX Block of Aloha St
2	5710	2015-01-28 09:30:00	12XX Block of E Pike St
3	2300	2015-02-21 08:24:21	15XX Block of Aurora St

typesID	Offense type	Offense code
1	tresspass	5700
2	burglary	5710
3	larceny-theft	2300

Database Implementation:

- there are many relational database software implementations:
 - commercial: Oracle, Microsoft SQL Server, IBM DB2
 - open source: MySQL, PostgreSQL
- Deployment
 - most databases are deployed on a server
 - can run locally for testing

Database Interface:

- databases are accessed via a *connection string*:
 - hostname, port, user, password
- one can connect through
 - command line
 - GUI apps (pg_admin, MySQL Workbench, DB Visualizer)
 - Python, R, etc.

In []:

In []:

In []:

In []: